

## 2. BUILDING A END TO END SPEECH RECOGNITION PIPE LINE SIGNAL PROCESSING ACOUSTIC MODELING & PERFORMANCE EVALUATION.

### **Program:**

```
import numpy as np
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, LSTM, Input
from sklearn.metrics import accuracy_score, classification_report
```

#### Step 1: Signal Processing (Simulated)

```
def preprocess_signal(signal, sample_rate=16000):
    """
    Simulate signal processing: Normalize and extract MFCC-like
    features.
    """
    # Normalize signal
    signal = signal / np.max(np.abs(signal))
    # Simulate feature extraction
    features = np.log1p(np.abs(np.fft.fft(signal))[:len(signal) // 2])
    return features
```

#### Step 2: Acoustic Modeling

```
def build_acoustic_model(input_dim, output_dim):
    """
    Build a simple LSTM-based acoustic model.
    """
    model = Sequential([
```

```

        Input(shape=(None, input_dim)),
        LSTM(128, return_sequences=False),
        Dense(output_dim, activation='softmax')
    ])

    model.compile(optimizer='adam',
loss='sparse_categorical_crossentropy', metrics=['accuracy'])

    return model

```

### # Step 3: Performance Evaluation

```
def evaluate_model(model, X_test, y_test):
```

```
    """
```

Evaluate the model on the test set.

```
    """
```

```

    predictions = model.predict(X_test)
    predicted_labels = np.argmax(predictions, axis=1)
    accuracy = accuracy_score(y_test, predicted_labels)
    report = classification_report(y_test, predicted_labels)
    return accuracy, report

```

### # Simulated Pipeline

```
if __name__ == "__main__":
```

```
    # Simulate raw audio signal (sine wave + noise)
```

```
    duration = 1.0 # seconds
```

```
    sample_rate = 16000
```

```

    time = np.linspace(0, duration, int(sample_rate * duration),
endpoint=False)

```

```

    signal = 0.5 * np.sin(2 * np.pi * 440 * time) + 0.1 *
np.random.randn(len(time))

# Step 1: Signal Processing
features = preprocess_signal(signal)

# Simulate Dataset
num_samples = 100
input_dim = features.shape[0]
X = np.random.rand(num_samples, 10, input_dim) # Simulated
features
y = np.random.randint(0, 5, size=(num_samples,)) # Simulated
labels (5 classes)

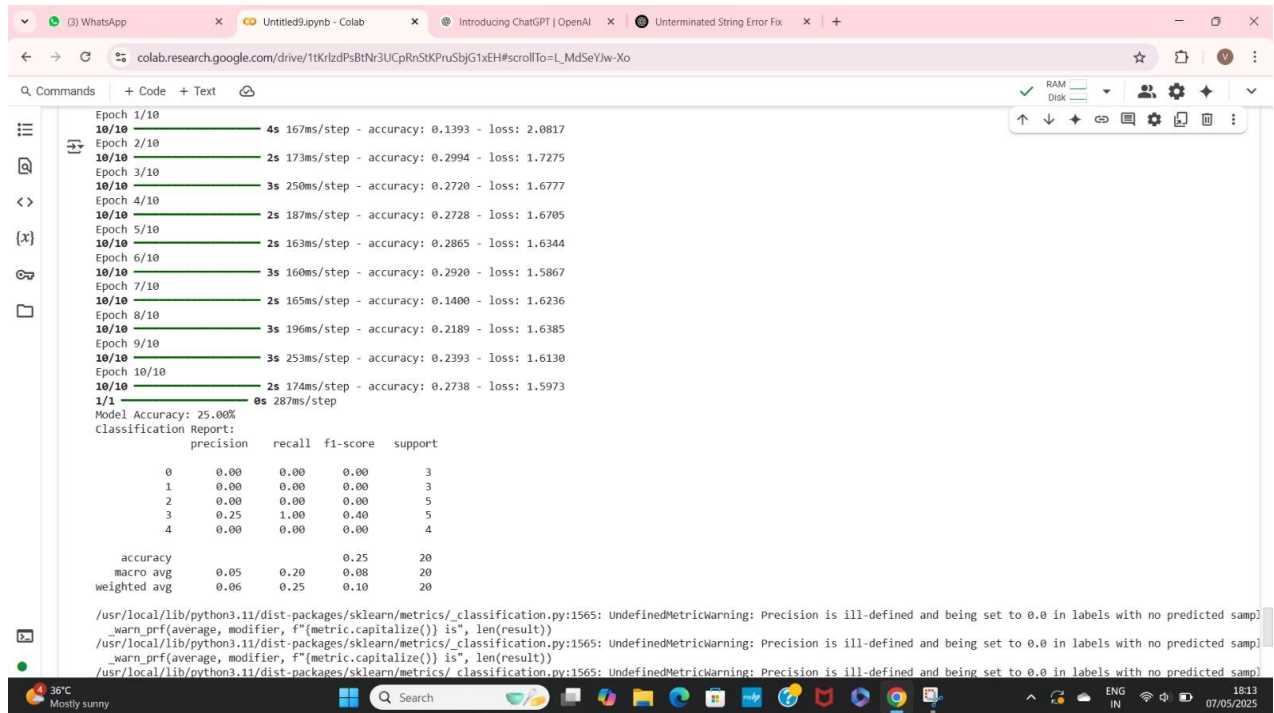
# Split into train and test sets
split_idx = int(0.8 * num_samples)
X_train, X_test = X[:split_idx], X[split_idx:]
y_train, y_test = y[:split_idx], y[split_idx:]

# Step 2: Acoustic Modeling
model = build_acoustic_model(input_dim, output_dim=5)
model.fit(X_train, y_train, epochs=10, batch_size=8, verbose=1)

# Step 3: Performance Evaluation
accuracy, report = evaluate_model(model, X_test, y_test)
print(f"Model Accuracy: {accuracy * 100:.2f}%")
print("Classification Report:\n", report)

```

output:



The screenshot shows a Google Colab notebook interface. The browser tabs at the top include 'WhatsApp', 'Untitled9.ipynb - Colab', 'Introducing ChatGPT | OpenAI', and 'Unterminated String Error Fix'. The address bar shows a Google Drive link. The notebook interface has a left sidebar with icons for file explorer, search, and code execution. The main area displays the output of a training process, showing progress bars for each of the 10 epochs. The final output is a 'Classification Report' with the following data:

```
Model Accuracy: 25.00%
Classification Report:
      precision    recall  f1-score   support

0         0.00      0.00      0.00         3
1         0.00      0.00      0.00         3
2         0.00      0.00      0.00         5
3         0.25      1.00      0.40         5
4         0.00      0.00      0.00         4

 accuracy          0.25      0.25      0.25      20
 macro avg          0.05      0.20      0.08      20
weighted avg          0.06      0.25      0.10      20
```

Below the report, there are three lines of warning messages from sklearn.metrics.\_classification.py:

```
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples
warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples
warn_prf(average, modifier, f'{metric.capitalize()} is', len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples
```

The bottom of the image shows a Windows taskbar with a search bar, system tray icons (including temperature at 36°C and date 07/05/2025), and the system clock showing 18:13.